

523454

Computer Network Programming

Multiplexing TCP connections

Dr. Parin Sornlertlamvanich,

parin.s@sut.ac.th

Multiplexing TCP connections

■ The socket APIs are blocking by default

- `accept()` to wait for an incoming connection, your program's execution is blocked until a new incoming connection
- `recv()` to read incoming data, your program's execution blocks until new data is available

■ Blocking I/O can be a significant problem

- Lab2's program needed to serve multiple clients

Multiplexing TCP connections (cont.)

- Then, imagine that one slow client connected to it
 - this slow client takes a minute before sending its first data
 - our server would simply be waiting on the `recv()` call to return
 - If other clients were trying to connect, they would have to wait it out
- Blocking also isn't usually acceptable on the client side either
 - In web browser, it has a tab feature where many whole web pages can be loaded in parallel
 - a technique for handling many separate connections simultaneously

(1) Polling non-blocking sockets

- To configure sockets to use a non-blocking operation
 - calling `fcntl()` with the `O_NONBLOCK` flag
- Once in non-blocking mode
 - a call to `recv()` with no data will return immediately
 - simply check each of its active sockets in turn
 - It would handle any socket that returned data and ignore any socket that didn't – called **Polling**

(1) Polling non-blocking sockets (cont.)

- Polling can be a waste of computer resources since most of the time
 - there will be no data to read
 - It also complicates the program (the programmer is required to manually track which sockets are active and which state, they are in)
 - Return values from **recv()** must also be handled differently than with blocking sockets

(2) Forking and multithreading

- To start a new thread or process for each connection
- In this case, blocking sockets are fine
 - they block only their servicing thread/process, and they do not block other threads/processes
- The `fork()` function splits the executing program into two separate processes

(2) Forking and multithreading (cont.)

```
while(1) {  
    socket_client = accept(socket_listen, &new_client, &new_client_length);  
    int pid = fork();  
    if (pid == 0) { //child process  
        close(socket_listen);  
        recv(socket_client, ...);  
        send(socket_client, ...);  
        close(socket_client);  
        exit(0);  
    }  
    //parent process  
    close(socket_client);  
}
```

- A multi-process TCP server may accept connections like this:
 - the program blocks on `accept()`
 - When a new connection is established, the program calls **`fork()`** to split into two processes
 - The child process, where `pid == 0`, only services this one connection
 - Using **`recv()`** freely without worrying about blocking
 - The parent process calls `close()` and returns to listening for more conns

(3) Synchronous multiplexing with select()

- `select()` <-> a set of sockets

- It tells us which ones are ready to be read
- It can also tell us which sockets are ready to write to

- `int select (int nfds, fd_set *readfds, fd_set *writefds, fd_set *exceptfds, struct timeval *timeout);`

(3) Synchronous multiplexing with select()

- `int select (maxDescPlus1, &readDescs, &writeDescs, &exceptionDescs, &timeout);`
 - `maxDescPlus1`: integer, hint of the maximum number of descriptors
 - `readDescs`: `fd_set`, checked for immediate input availability
 - `writeDescs`: `fd_set`, checked for the ability to immediately write data
 - `exceptionDescs`: `fd_set`, checked for pending exceptions
 - `timeout`: `struct timeval`, how long it blocks (NULL → forever)
- **returns** the total number of ready descriptors, -1 in case of error
- **changes** the descriptor lists so that only the corresponding positions are set

```
int FD_ZERO (fd_set *descriptorVector); /* removes all descriptors from vector */
int FD_CLR (int descriptor, fd_set *descriptorVector); /* remove descriptor from vector */
int FD_SET (int descriptor, fd_set *descriptorVector); /* add descriptor to vector */
int FD_ISSET (int descriptor, fd_set *descriptorVector); /* vector membership check */
```

```
struct timeval {
    time_t tv_sec; /* seconds */
    time_t tv_usec; /* microseconds */
};
```

(3) Synchronous multiplexing with select()

- Before calling `select()`, we must first add our sockets into an `fd_set`
- If we have three sockets, `socket_listen`, `socket_a`, and `socket_b`
- To add them to an `fd_set`

```
fd_set our_sockets;  
FD_ZERO(&our_sockets);  
FD_SET(socket_listen, &our_sockets);  
FD_SET(socket_a, &our_sockets);  
FD_SET(socket_b, &our_sockets);
```

It is important to zero-out the `fd_set` using `FD_ZERO()` before use

(3) Synchronous multiplexing with select()

- `select()` also requires that we pass a number that's larger than the largest socket descriptor

```
int max_socket;  
max_socket = socket_listen;  
if (socket_a > max_socket) max_socket = socket_a;  
if (socket_b > max_socket) max_socket = socket_b;
```

- When we call `select()`, it modifies our `fd_set` of sockets to indicate which sockets are ready
 - Therefore, we want to copy our socket set before calling it

```
fd_set copy;  
copy = our_sockets;  
  
select(max_socket+1, &copy, 0, 0, 0);
```

(3) Synchronous multiplexing with select()

- This call blocks until at least one of the sockets is ready to be **read** from
- When select() returns
 - **copy** is modified so that it only contains the sockets that are ready to be read from
 - To check which sockets are still in copy using **FD_ISSET()**

```
if (FD_ISSET(socket_listen, &copy)) {  
    //socket_listen has a new connection  
    accept(socket_listen...  
}  
if (FD_ISSET(socket_a, &copy)) {  
    //socket_a is ready to be read from  
    recv(socket_a...  
}  
if (FD_ISSET(socket_b, &copy)) {  
    //socket_b is ready to be read from  
    recv(socket_b...  
}
```

(3) Synchronous multiplexing with `select()`

- If we wanted to monitor an `fd_set` for writability instead of readability
 - pass our `fd_set` as the third argument to `select()`
- We can monitor a set of sockets for exceptions by passing it as
 - the fourth argument to `select()`

select() timeout

■ The last argument allows us to specify a timeout

- `tv_sec` holds the number of seconds, and `tv_usec` holds the number of microseconds

```
struct timeval {  
    long tv_sec;  
    long tv_usec;  
}
```

■ If we want select() to wait a maximum of 1.5 seconds

- `select()` returns after a socket in `fd_set copy` is ready to read or after 1.5 seconds has elapsed, whichever is sooner

```
struct timeval timeout;  
timeout.tv_sec = 1;  
timeout.tv_usec = 500000;  
select(max_socket+1, &copy, 0, 0, &timeout);
```

(3) Synchronous multiplexing with select()

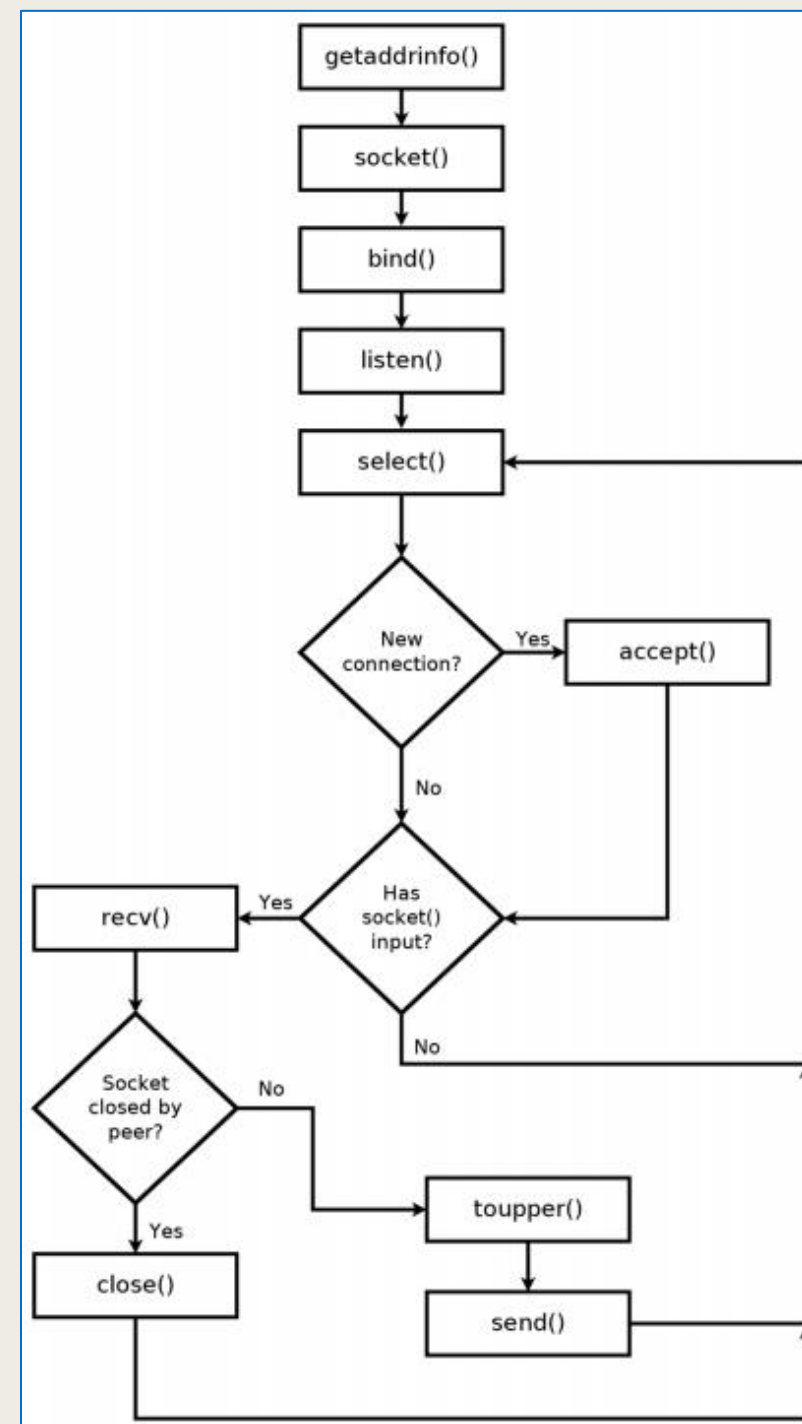
- select() can also be used to monitor for writeable sockets, and sockets with exceptions
 - sockets where we could call **send()** without blocking

```
select(max_sockets+1, &ready_to_read, &ready_to_write,  
&excepted, &timeout);
```

- On success, select() itself returns the number of socket descriptors
 - The return value is zero if it timed out before
 - returns -1 to indicate an error

Example: TCP server

- TCP server that converts strings into uppercase
- If a client connects and sends Hello
 - The server will send **HELLO** back
- To call `select()`, which alerts us if a new connection is available
 - a new connection is waiting do we call `accept()`
- When data is received by `recv()`, we run it through `toupper()`
 - return it to the client using `send()`



■ getaddrinfo() -> socket() -> bind() -> listen()

```
struct addrinfo *bind_address;
getaddrinfo(0, "7777", &hints, &bind_address);

printf("Creating socket...\n");
int socket_listen;
socket_listen = socket(bind_address->ai_family,
                      bind_address->ai_socktype, bind_address->ai_protocol);

printf("Binding socket to local address...\n");
if (bind(socket_listen,
        bind_address->ai_addr, bind_address->ai_addrlen)) {
    fprintf(stderr, "bind() failed. (%d)\n", GETSOCKETERRNO());
    return 1;
}
freeaddrinfo(bind_address);

printf("Listening...\n");
if (listen(socket_listen, 10) < 0) {
    fprintf(stderr, "listen() failed. (%d)\n", GETSOCKETERRNO());
    return 1;
}
```

Example: TCP server

- `getaddrinfo()` -> `socket()` -> `bind()` -> `listen()`
- To define an `fd_set` structure that stores all of the active sockets
 - For now, we add only our **listening socket** -> **`max_socket`** (it's the only socket)

```
fd_set master;  
FD_ZERO(&master);  
FD_SET(socket_listen, &master);  
SOCKET max_socket = socket_listen;
```

Example: TCP server

■ To add new connections to master

- To enter the **main loop**, and set up our call to **select()**

```
printf("Waiting for connections...\n");

while(1) {
    fd_set reads;
    reads = master;
    if (select(max_socket+1, &reads, 0, 0, 0) < 0) {
        fprintf(stderr, "select() failed. (%d)\n", GETSOCKETERRNO());
        return 1;
    }
    ...
}
```

- This works by first copying our fd_set **master** into **reads**.
- Recall that **select()** modifies the set given to it. If we didn't copy master, we would lose its data
- Timeout value of 0 -> it doesn't return until a socket in the master set is ready to be read from

Example: TCP server

- Whether it was flagged by `select()` as being ready. If a socket, `X`, was flagged by `select()`
 - then `FD_ISSET(X, &reads)` is true
- Socket descriptors are positive integers, so we can try every possible socket descriptor up to `max_socket`

```
int i;  
for(i = 1; i <= max_socket; ++i) {  
    if (FD_ISSET(i, &reads)) {  
        //Handle socket – Next slide  
    }  
}
```

Example: TCP server

- `FD_ISSET()` is only true for sockets that are ready to be read
- In the case of `socket_listen`, this means that a new connection is ready to be established with `accept()`
 - We should first determine whether the current socket is the listening one or not. If it is, we call `accept()`
- For all other sockets, it means that data is ready to be read with `recv()`

```

if (i == socket_listen) {
    struct sockaddr_storage client_address;
    socklen_t client_len = sizeof(client_address);
    int socket_client = accept(socket_listen,
        (struct sockaddr*) &client_address,
        &client_len);
    if (socket_client == -1) {
        fprintf(stderr, "accept() failed. (%d)\n",
            GETSOCKETERRNO());
        return 1;
    }

    FD_SET(socket_client, &master);
    if (socket_client > max_socket)
        max_socket = socket_client;

    char address_buffer[100];
    getnameinfo((struct sockaddr*)&client_address,
        client_len,
        address_buffer, sizeof(address_buffer), 0, 0,
        NI_NUMERICHOST);
    printf("New connection from %s\n", address_buffer);
} else {
    ...

```

■ If the socket `i` is not `socket_listen`

- then it is instead a request for an established connection (**existing connection**)
- to read it with **`recv()`**, convert it into **uppercase**

```
} else {  
    char read[1024];  
    int bytes_received = recv(i, read, 1024, 0);  
    if (bytes_received < 1) {  
        FD_CLR(i, &master);  
        close(i);  
        continue;  
    }  
    int j;  
    for (j = 0; j < bytes_received; ++j)  
        read[j] = toupper(read[j]);  
    send(i, read, bytes_received, 0);  
}
```

Example: TCP server

```
        } //if FD_ISSET
    } //for i to max_socket
} //while(1)

printf("Closing listening socket...\n");
close(socket_listen);

printf("Finished.\n");
return 0;
}
```


Example: TCP server

```
(kali㉿kali)-[~/lab_netPro/my_lab/lab3/temp]
$ ./serv
Configuring local address ...
Creating socket ...
Binding socket to local address ...
Listening ...
Waiting for connections ...
New connection from 127.0.0.1
New connection from 127.0.0.1
New connection from 127.0.0.1
█
```

```
(kali㉿kali)-[~/lab_netPro/my_lab/lab3/temp]
$ ./client 127.0.0.1 7777
Configuring remote address ...
Remote address is: 127.0.0.1 7777
Creating socket ...
Connecting ...
Connected.
To send data, enter text followed by enter.
CPE2
Sending: CPE2
Sent 5 bytes.
Received (5 bytes): CPE2
█
```

```
(kali㉿kali)-[~/lab_netPro/my_lab/lab3/temp]
$ ./client 127.0.0.1 7777
Configuring remote address ...
Remote address is: 127.0.0.1 7777
Creating socket ...
Connecting ...
Connected.
To send data, enter text followed by enter.
CPE1
Sending: CPE1
Sent 5 bytes.
Received (5 bytes): CPE1
```

```
(kali㉿kali)-[~/lab_netPro/my_lab/lab3/temp]
$ ./client 127.0.0.1 7777
Configuring remote address ...
Remote address is: 127.0.0.1 7777
Creating socket ...
Connecting ...
Connected.
To send data, enter text followed by enter.
CPE3
Sending: CPE3
Sent 5 bytes.
Received (5 bytes): CPE3
SUT3
Sending: SUT3
Sent 5 bytes.
Received (5 bytes): SUT3
```